

DSC 40B

Lecture 15 : Hashing

Warmup

Exercise

- How fast can we query/insert with these data structures?

	Query	Insert
Unsorted linked list	$\Theta(?)$	$\Theta()$
Unsorted array	$\Theta()$	$\Theta()$
Sorted array	$\Theta()$	$\Theta()$
Balanced BST	$\Theta()$	$\Theta()$
Unbalanced BST	$\Theta()$	$\Theta()$

Exercise

- How fast can we query/insert with these data structures?

	Query	Insert
Unsorted linked list	$\Theta(n)$	$\Theta(?)$
Unsorted array	$\Theta(?)$	$\Theta(?)$
Sorted array	$\Theta(?)$	$\Theta(?)$
Balanced BST	$\Theta(?)$	$\Theta(?)$
Unbalanced BST	$\Theta(?)$	$\Theta(?)$

Exercise

- How fast can we query/insert with these data structures?

	Query	Insert
Unsorted linked list	$\Theta(n)$	$\Theta(1)$
Unsorted array	$\Theta(?)$	$\Theta(?)$
Sorted array	$\Theta(?)$	$\Theta(?)$
Balanced BST	$\Theta(?)$	$\Theta(?)$
Unbalanced BST	$\Theta(?)$	$\Theta(?)$

Exercise

- How fast can we query/insert with these data structures?

	Query	Insert
Unsorted linked list	$\Theta(n)$	$\Theta(1)$
Unsorted array	$\Theta(n)$	$\Theta(?)$
Sorted array	$\Theta(?)$	$\Theta(?)$
Balanced BST	$\Theta(?)$	$\Theta(?)$
Unbalanced BST	$\Theta(?)$	$\Theta(?)$

Exercise

- How fast can we query/insert with these data structures?

	Query	Insert
Unsorted linked list	$\Theta(n)$	$\Theta(1)$
Unsorted array	$\Theta(n)$	$\Theta(n)$
Sorted array	$\Theta(?)$	$\Theta(?)$
Balanced BST	$\Theta(?)$	$\Theta(?)$
Unbalanced BST	$\Theta(?)$	$\Theta(?)$

Exercise

- How fast can we query/insert with these data structures?

	Query	Insert
Unsorted linked list	$\Theta(n)$	$\Theta(1)$
Unsorted array	$\Theta(n)$	$\Theta(n)$
Sorted array	$\Theta(\log n)$	$\Theta(?)$
Balanced BST	$\Theta(?)$	$\Theta(?)$
Unbalanced BST	$\Theta(?)$	$\Theta(?)$

Exercise

- How fast can we query/insert with these data structures?

	Query	Insert
Unsorted linked list	$\Theta(n)$	$\Theta(1)$
Unsorted array	$\Theta(n)$	$\Theta(n)$
Sorted array	$\Theta(\log n)$	$\Theta(n)$
Balanced BST	$\Theta(?)$	$\Theta(?)$
Unbalanced BST	$\Theta(?)$	$\Theta(?)$

Exercise

- How fast can we query/insert with these data structures?

	Query	Insert
Unsorted linked list	$\Theta(n)$	$\Theta(1)$
Unsorted array	$\Theta(n)$	$\Theta(n)$
Sorted array	$\Theta(\log n)$	$\Theta(n)$
Balanced BST	$\Theta(\log n)$	$\Theta(?)$
Unbalanced BST	$\Theta(?)$	$\Theta(?)$

Exercise

- How fast can we query/insert with these data structures?

	Query	Insert
Unsorted linked list	$\Theta(n)$	$\Theta(1)$
Unsorted array	$\Theta(n)$	$\Theta(n)$
Sorted array	$\Theta(\log n)$	$\Theta(n)$
Balanced BST	$\Theta(\log n)$	$\Theta(\log n)$
Unbalanced BST	$\Theta(?)$	$\Theta(?)$

Exercise

- How fast can we query/insert with these data structures?

	Query	Insert
Unsorted linked list	$\Theta(n)$	$\Theta(1)$
Unsorted array	$\Theta(n)$	$\Theta(n)$
Sorted array	$\Theta(\log n)$	$\Theta(n)$
Balanced BST	$\Theta(\log n)$	$\Theta(\log n)$
Unbalanced BST	$\Theta(n)$	$\Theta(?)$

Exercise

- How fast can we query/insert with these data structures?

	Query	Insert
Unsorted linked list	$\Theta(n)$	$\Theta(1)$
Unsorted array	$\Theta(n)$	$\Theta(n)$
Sorted array	$\Theta(\log n)$	$\Theta(n)$
Balanced BST	$\Theta(\log n)$	$\Theta(\log n)$
Unbalanced BST	$\Theta(n)$	$\Theta(n)$

Direct Address Tables

Counting Frequencies

- How many times does each age appear?

PID	Name	Age
A1843	Wan	24
A8293	Deveron	22
A9821	Vinod	41
A8172	Aleix	17
A2882	Kayden	4
A1829	Raghu	51
A9772	Cui	48
⋮	⋮	⋮

Exercise

What data structure would you use to store the age counts?

Direct Address Tables

- **Idea**: keep an **array** of length, say, 125.
- Initialize to zero.

- | | | | | | | | |
|---|---|---|---|-----|---|---|-----|
| 0 | 0 | 0 | 0 | ... | 0 | 0 | 0 |
| 0 | 1 | 2 | 3 | ... | 5 | 6 | 124 |

Direct Address Tables

- **Idea:** keep an **array** `arr` of length, say, 125.
- Initialize to zero.

0	0	0	0	...	0	0	0
0	1	2	3	...	5	6	124

- If we see age x , increment `arr[x]` by one. Say, we see 2:
 - `arr[2] += arr[2] + 1`

0	0	1	0	...	0	0	0
0	1	2	3	...	5	6	124

Building the Table

```
# loading the table  
table = np.zeros(125)  
  
for age in ages:  
    table[age] += 1
```

Building the Table

```
# loading the table
table = np.zeros(125)

for age in ages:
    table[age] += 1
```

Time complexity if there are n people?

A: Constant

B: $\log n$

C: n

D: n^2

Query

query: how many people are 55?

```
print(table[55])
```

- Time complexity if there are n people?

Time complexity if there are n people?

A: Constant

B: $\log n$

C: n

D: n^2

Query

query: how many people are 55?

```
print(table[55])
```

- Time complexity if there are n people?
 - $\Theta(1)$

Counting Names

PID	Name	Age
A1843	Wan	24
A8293	Deveron	22
A9821	Vinod	41
A8172	Aleix	17
A2882	Kayden	4
A1829	Raghu	51
A9772	Cui	48
:	:	:

Downsides

- Direct Access Tables are **fast**.
- What are the downsides of DATs?
- Could we use a DAT to store:
 - zip codes?
 - phone numbers?
 - credit card numbers?
 - names?

Downsides

- Things being stored must be integers, or convertible to integers.
 - why? valid array indices
- Must come from a small range of possibilities
 - why? memory usage.
 - example: phone numbers

Running time for Direct Hashing or DAT

```
class StudentDataBase:
    allStudents = np.zeros(4294967268)

    def add (s):
        index = s.studentId
        allStudents[index] = s

    def get (s):
        index = s.studentId
        return allStudent[index]

    def remove(s):
        index = s.studentId
        allStudents[index] = None
```

Space complexity is horrible :(

- Allocating 4 billion entries for a UCSD student database is extremely wasteful.
 - The universe of keys is far larger than the number we expect to ever want to store.
- What if we decrease the size of the table from 4 billion entries to storing, say, just 100,000?
- 100,000 is still far higher than the actual number of UCSD students, so there should be plenty of space.

Hash Tables

- **Insight:** anything can be “converted” to an integer through **hashing**.
- But not *uniquely*!
- Hash tables have many of the same advantages as DATs, but work more generally.

Hashing

Hashing

- One of the most important ideas in CS.
- **Tons** of uses:
 - Verifying message integrity.
 - Fast queries on a large data set.
 - Identify if file has changed in version control.

Hash Function

- A **hash function** takes a (large) *object* and returns a (smaller) “fingerprint” of that object.
- Usually the fingerprint is a number, guaranteed to be in some range.

How?

- Looking at certain bits, combining them in ways that **look** random (but aren't!)

Hash Function Properties

- Hashing same thing **twice** returns the **same** hash.
- Unlikely that different things have same fingerprint.
 - But not impossible!

Collisions

- Hash functions map objects to numbers in a defined range.
 - **Example**: given image, return number in $[0, 1, 2, \dots, 1024]$
- There will be two images with the same hash.
 - **Pigeonhole principle**: if there are n pigeons, $< n$ holes, there will a hole with ≥ 2 pigeons.
- **Collision**: two objects have the **same** hash

“Good” Hash Functions

- A **good hash** function tries to minimize collisions.

Hashing in Python

The **hash** function computes a hash.

```
>>> hash("This is a test")  
-670458579957477203
```

```
>>> hash("This is a test")  
-670458579957477203
```

```
>>> hash("This is a test!")  
1860306055874153109
```


Example

- MD5 is a **cryptographic** hash function.
 - Hard to “reverse engineer” input from hash.
- Returns a really large number in `hex`.

a741d8524a853cf83ca21eabf8cea190

- Used to “fingerprint” whole files.

Example

```
> echo "My name is Marina" | md5
```

```
60b5504d3410d0f97796c95829adebfd
```

```
> echo "My name is Marina" | md5
```

```
60b5504d3410d0f97796c95829adebfd
```

```
> echo "My name is marina" | md5
```

```
4a3c494e411e94d2be7c32ac7461208f
```

Example

```
>>> md5 finalA.pdf
```

```
MD5 (finalA.pdf) = ff5bcc7dcc4815aa221774b1507c42c5
```

Why? Useful for security reasons

- I release a piece of software.
- I host it on Google Drive.
- Someone (Google, US Gov., etc.) decides to insert extra code into software to spy on users.
- You have no way of knowing.

Another Use: De-duplication

- Building a massive training set of images.
- Given a new image, is it already in my collection?
- Don't need to compare images pixel-by-pixel!
- Instead, compare **hashes**.

Hashing for Data Scientists

- Don't need to know much about *how* the hash function works.
- But should know how they are used.

Hash Tables

Membership Queries

- **Given:** a collection of n numbers and a target t .
- **Find:** determine if t is in the collection.

Goal

- **D**irect **A**ccess **T**ables are fast, but won't work for things that aren't numbers in a small range.
- **Idea**: hash objects to numbers in a small range, use a DAT.
- But must **deal with collisions**.

Hash Tables

- Pick a **table size** m .
 - Usually $m \approx$ number of things you'll be storing.
- Create hash function to turn input into a number in $\{0, 1, \dots, m - 1\}$.
- Create array with m bins.


Example

```
hash('hello') == 3  
hash('data') == 0  
hash('science') == 4
```

0 1 2 3 4 ... $m - 1$


Example

```
hash('hello') == 3  
hash('data') == 0  
hash('science') == 4
```



Example

```
hash('hello') == 3  
hash('data') == 0  
hash('science') == 4
```



Example

```
hash('hello') == 3  
hash('data') == 0  
hash('science') == 4
```

<u>"data"</u>			<u>"hello"</u>	<u>"science"</u>	...	<u></u>
0	1	2	3	4		$m - 1$

Collisions

- The **universe** is the set of all *possible* inputs.
- This is usually **much larger** than m (even infinite).
- **Not possible** to assign each input to a unique bin.
- If $\text{hash}(a) == \text{hash}(b)$, there is a **collision**.

Example

```
hash('hello') == 3
```

```
hash('data') == 0
```

```
hash('science') == 3
```


0 1 2 3 4 ... $m - 1$

Example

```
hash('hello') == 3
```

```
hash('data') == 0
```

```
hash('science') == 3
```




Example

```
hash('hello') == 3
```

```
hash('data') == 0
```

```
hash('science') == 3
```




Example

```
hash('hello') == 3
```

```
hash('data') == 0
```

```
hash('science') == 3
```




Example

```
hash('hello') == 3
```

```
hash('data') == 0
```

```
hash('science') == 3
```




Example

```
hash('hello') == 3
```


```
hash('data') == 0
```

```
hash('science') == 3
```



Chaining

- Collisions stored in same bin, in linked list.
- **Query:** *Hash* to find bin, then *linear* search.



The Idea

- A **good** hash function will utilize all bins evenly.
 - Looks like **uniform** random distribution.
- If $m \approx n$, then **only a few** elements in each bin.
- As we add more elements, we need to add bins.

Expected Case

- n elements.
- m bins.
- Assume elements placed **randomly** (but *deterministically*) in bins.
- **Expected** size of linked list inside the bin: n/m .

Analysis

- **Query:**
 - $\Theta(1)$ to find bin (hashing step)
 - $\Theta(n/m)$ for a linear search in a list.
 - **Total:** $\Theta(1 + n/m)$. ($m \approx n$)
 - We *usually* guarantee $m = \Theta(n)$
 - **Expected time:** $\Theta(1)$.

Worst Case

- Everything **hashed to same bin**.
 - Really unlikely!
- **Query:**
 - $\Theta(1)$ to find bin
 - $\Theta(n)$ for linear search
 - Total: $\Theta(n)$

Exercise

What is the worst case time complexity of **inserting** an element into a hash table that uses chaining with linked lists?

Growing the Hash Table

- Insertions take $\Theta(1)$ **unless** the hash table needs to *grow*.
- We need to ensure that $m \leq c \cdot n$.
 - Otherwise, **too many** collisions.
- If we add a bunch of elements, we'll need to **increase** m .
- Increasing m mean allocating a new array, $\Theta(m) = \Theta(n)$ time.

Main Idea

Hash tables support **constant** (*expected*) time insertion and membership queries.

Dictionaries

- Hash tables can also be used to store (key, value) pairs.
- Often called `dictionaries` or `associative arrays`.

Hashing in Python

- dict and set implement hash tables.
- Querying is done using **in**:

make a set

```
>>> L = {3, 6, -2, 1, 7, 12}
```

```
>>> 1 in L          # Theta(1)
```

```
True
```

```
>>> 8 in L          # Theta(1)
```

```
False
```

Hashing in Python

- Querying is done using **in**:

make a list

```
>>> L = [3, 6, -2, 1, 7, 12]
```

```
>>> 1 in L      # Theta(?)
```

True

```
>>> 8 in L      # Theta(?)
```

False

Hashing in Python

- Querying is done using **in**:

make a list


```
>>> L = [3, 6, -2, 1, 7, 12]
```

```
>>> 1 in L      # Theta(n)
```

True

```
>>> 8 in L      # Theta(n)
```

False



Fast Algorithms with Hash Tables

Faster Algorithms

- Hashing is a **super** common trick.
- The “best” solution to interview problems **often** involves hashing.

Example 1: The Movie Problem

- You're on a flight that will last D minutes.
- You want to pick two movies to watch.
- Find two whose durations sum to **exactly** D .

Recall: Previous Solutions

- Brute force: $\Theta(n^2)$.
- Sort, use sorted structure: $\Theta(n \log n) + \Theta(n)$.
- Theoretical lower bound: $\Omega(n)$?
- Can we speed this up with **hash tables**?

Idea

- To use hash tables, we want to frame problem as a **membership** query.

Example

- Suppose flight is 360 (D) minutes long.
- Suppose first movie is fixed: 120 minutes.
- Is there a movie lasting $(360 - 120) = 240$ minutes?

```
def optimize_entertainment_hash(times, D):  
    hash_table = dict()  
    for i, time in enumerate(times):  
        hash_table[time] = i  
  
    for i, time in enumerate(times):  
        target = D - time  
        if target in hash_table:  
            return i, hash_table[target]
```


Example 2: Anagrams

- Two strings w_1 and w_2 are **anagrams** if the letters of w_1 one can be *permuted* to make w_2 .

Examples of anagrams

- abcd / dbca
- listen / silent
- sandiego / doginsea

Problem

- Given a collection of n strings, determine if any two of them are anagrams.
- Design an efficient algorithm for solving this problem. What is its time complexity?

Solution

- Let's turn this into a *membership* query.
- **Trick:** two strings are anagrams if

$$\text{sorted}(w_1) == \text{sorted}(w_2)$$

```
def any_anagrams(words):  
    seen = set()  
    for word in words:  
        w = sorted(word)  
        if w in seen:  
            return True  
        else:  
            seen.add(w)
```

What is the worst time complexity?

A: Constant

B: n

C: $n \log n$

D: n^2

Hashing **Downsides**

- Problem must involve **membership** query.

Example: The Movie Problem

- You're on a flight that will last D minutes.
- You want to pick two movies to watch.
- Find two whose added durations is **closest** to D .

Hashing **Downsides**

- **No cache locality**: similar items map to different bins.
- There is no way to quickly query entry closest to given input.

Example: Number of Elements

- Given a collection of n numbers and two endpoints, a and b , determine how many of the numbers are contained in $[a, b]$.
- Not a membership query.
- **Idea:** **sort** and use *modified* binary search.

Hash Table Drawbacks

Hashing **Downsides**

- No locality: similar items map to different bins.
- But we often want similar items at the same time.
- Results in many **cache misses**, **slow**.
- Memory overhead.

Hash Tables vs. BSTs

- **Hash Table:** $\Theta(1)$ insertion, query (expected time).
- **BST:** $\Theta(\log n)$ insertion, query (if *balanced*).
- Why ever use a BST?

Hash Tables vs. BSTs

- Hash tables keep items in arbitrary order.
- **Example:** how many elements are in the interval $[3, 23]$?
- **Example:** what is the min/max/median?
- BSTs win when order is important.



Thank you!

Do you have any questions?

CampusWire!