


DSC 40B

***Lecture 29 : Minimum
Spanning Trees.
Kruskal's algorithm***

Minimum Spanning Trees

Last time: Minimum Spanning Tree


- The **minimum spanning** tree problem is as follows:
 - **Given:** A weighted, undirected graph $G = (V, E, \omega)$.
 - **Compute:** a spanning tree of G with minimum cost (i.e., minimum total edge weight).
- For a given graph, the MST may not be unique.



Cost:

$$1+2+3+4+5+7+8=30$$

Prim's Algorithm, Informally



- Start by picking any node to add to “tree”, T .
- While T is not a spanning tree, greedily add **lightest** edge from a node in T to a node not in T .
 - “Lightest” = edge with the smallest weight.
- **Is this guaranteed to work?** Yes, as we’ll see.


Last time: Building MSTs

- How do we build a MST efficiently?
- We'll adopt a **greedy** approach.
 - Build a tree edge-by-edge.
 - At every step, doing what looks best at the moment.
- This strategy isn't guaranteed to work in all of life's situations, but it works for building MSTs.

Last time: Two Greedy Approaches


- We'll look at two greedy algorithms:
 - **Last time:** Prim's Algorithm
 - **Today: *Kruskal's Algorithm***
- Differ in the order in which edges are added to tree.
- Also differ in time complexity.

Prim's Algorithm, Informally




- Start by picking **any** node to add to “tree”, T .
- While T is **not** a spanning tree, greedily add **lightest** edge from a node **in** T to a node **not in** T .
 - “*Lightest*” = edge with the smallest weight.

Kruskal's Algorithm, Informally




- Start with **empty** forest: $T = (V, E_{\text{mst}})$, where $E_{\text{mst}} = \emptyset$.
- Loop through **edges** in *increasing* order of weight.
 - If edge **does not** create a cycle in T , add it to T .
 - If T is a **spanning tree**, break.

Kruskal's Algorithm, Informally




- Start with **empty** forest: $T = (V, E_{\text{mst}})$, where $E_{\text{mst}} = \emptyset$.
- Loop through **edges** in *increasing* order of weight.
 - If edge **does not** create a cycle in T , add it to T .
 - If T is a **spanning tree**, break.

Kruskal's Algorithm, Informally




- Start with **empty** forest: $T = (V, E_{\text{mst}})$, where $E_{\text{mst}} = \emptyset$.
- Loop through **edges** in *increasing* order of weight.
 - If edge **does not** create a cycle in T , add it to T .
 - If T is a **spanning tree**, break.

Kruskal's Algorithm, Informally




- Start with **empty** forest: $T = (V, E_{\text{mst}})$, where $E_{\text{mst}} = \emptyset$.
- Loop through **edges** in *increasing* order of weight.
 - If edge **does not** create a cycle in T , add it to T .
 - If T is a **spanning tree**, break.

Kruskal's Algorithm, Informally




- Start with **empty** forest: $T = (V, E_{\text{mst}})$, where $E_{\text{mst}} = \emptyset$.
- Loop through **edges** in *increasing* order of weight.
 - If edge **does not** create a cycle in T , add it to T .
 - If T is a **spanning tree**, break.

Kruskal's Algorithm, Informally




- Start with **empty** forest: $T = (V, E_{\text{mst}})$, where $E_{\text{mst}} = \emptyset$.
- Loop through **edges** in *increasing* order of weight.
 - If edge **does not** create a cycle in T , add it to T .
 - If T is a **spanning tree**, break.

Kruskal's Algorithm, Informally




- Start with **empty** forest: $T = (V, E_{\text{mst}})$, where $E_{\text{mst}} = \emptyset$.
- Loop through **edges** in *increasing* order of weight.
 - If edge **does not** create a cycle in T , add it to T .
 - If T is a **spanning tree**, break.

Kruskal's Algorithm, Informally




- Start with **empty** forest: $T = (V, E_{\text{mst}})$, where $E_{\text{mst}} = \emptyset$.
- Loop through **edges** in *increasing* order of weight.
 - If edge **does not** create a cycle in T , add it to T .
 - If T is a **spanning tree**, break.

Kruskal's Algorithm, Informally



- Start with **empty** forest: $T = (V, E_{\text{mst}})$, where $E_{\text{mst}} = \emptyset$.
- Loop through **edges** in *increasing* order of weight.
 - If edge **does not** create a cycle in T , add it to T .
 - If T is a **spanning tree**, break.

Kruskal's Algorithm, Informally



- Start with **empty** forest: $T = (V, E_{\text{mst}})$, where $E_{\text{mst}} = \emptyset$.
- Loop through **edges** in *increasing* order of weight.
 - If edge **does not** create a cycle in T , add it to T .
 - If T is a **spanning tree**, break.

Being Greedy

- **Prim**: add the **node** with smallest estimated cost and update neighbors.
 - Works locally, “grows” a connected tree.
- **Kruskal**: add the **edge** with smallest weight.
 - As long as it doesn’t make a cycle.
 - Edge can be anywhere in graph.

Kruskal's Algorithm (Pseudocode)

```
def kruskal(graph, weights):  
    mst = UndirectedGraph()  
  
    # sort edges in ascending order by weight  
    sorted_edges = sorted(graph.edges, key=weights.get)  
  
    for (u, v) in sorted_edges:  
        # if u and v are not already connected  
        if ...:  
            mst.add_edge(u, v)  
  
        # (optional) if mst is now a spanning tree, break  
        if len(mst.edges) == len(graph.nodes) - 1:  
            break  
  
    return mst
```

Kruskal's Algorithm (Pseudocode)

```
def kruskal(graph, weights):  
    mst = UndirectedGraph()  
  
    # sort edges in ascending order by weight  
    sorted_edges = sorted(graph.edges, key=weights.get)  
  
    for (u, v) in sorted_edges:  
        # if u and v are not already connected  
        if ...:  
            mst.add_edge(u, v)  
  
        # (optional) if mst is now a spanning tree, break  
        if len(mst.edges) == len(graph.nodes) - 1:  
            break  
  
    return mst
```

To check if u and v are connected

A: BFS

B: DFS

C: None

D: More than one

Checking for Connectivity

- Each iteration: check if u and v are connected in $T = (V, E_{\text{mst}})$.
- We *could* do a DFS/BFS on each iteration but ...
 - $\Theta(V + E_{\text{mst}}) = \Theta(V)$ each time.
 - **Expensive!**
- **Remember:**
 - If you're computing something **once**, use a *fast algorithm*.
 - If you're computing it **repeatedly**, consider a *data structure*.

Disjoint Set Forests

- Represent a collection of disjoint sets.

$\{\{1, 5, 6\}, \{2, 3\}, \{0\}, \{4\}\}$

- `.union(x, y)`: Union the sets containing x and y .
- `.in_same_set(x, y)`: Return **True/False** if x and y are in the same set*.

**Usually implemented as a `.find(x)` method returning representative of set containing x .*

Example

```
>>> # create a DSF with {{0}, {1}, {2}, {3}, {4}, {5}}
>>> dsf = DisjointSetForest([0, 1, 2, 3, 4, 5])
>>> dsf.union(0, 3)
>>> dsf.union(1, 4)
>>> dsf.union(3, 1)
>>> dsf.union(2, 5)
>>> # dsf now represents {{0, 1, 3, 4}, {2, 5}}
>>> dsf.in_same_set(0, 3)
True
>>> dsf.in_same_set(0, 2)
False
```

Disjoint Set Forests

- Operations take $\Theta(\alpha(n))$ time, where n is number of objects in collection.
- $\alpha(n)$ is the **inverse Ackermann function**.
- It grows **very, very** slowly.
- *Essentially* constant time.

Disjoint Set Forests

- Can be used to keep track of Connected Components of a **dynamic graph**.
- Nodes of Connected Components are disjoint sets.
 - Add an edge (u, v) : `.union(u, v)`
 - Check if u and v are connected: `.in_same_set(u, v)`
- To check if u, v are already connected:
 - BFS/DFS: $\Theta(V)$ each time.
 - DSF: $\Theta(\alpha(V))$ each time (essentially $\Theta(1)$).

```
def kruskal(graph, weights):  
    mst = UndirectedGraph()  
  
    # place each node in its own disjoint set  
    components = DisjointSetForest(graph.nodes)  
  
    # sort edges in ascending order by weight  
    sorted_edges = sorted(graph.edges, key=weights)  
  
    for (u, v) in sorted_edges:  
        if not components.in_same_set(u, v):  
            mst.add_edge(u, v)  
            components.union(u, v)  
  
            # (optional) if mst is now a spanning tree, break  
            if len(mst.edges) == len(graph.nodes) - 1:  
                break  
  
    return mst
```

Time Complexity

$\Theta(V)$

```
def kruskal(graph, weights):  
    mst = UndirectedGraph()
```

```
    # place each node in its own disjoint set  
    components = DisjointSetForest(graph.nodes)
```

```
    # sort edges in ascending order by weight  
    sorted_edges = sorted(graph.edges, key=weights)
```

```
    for (u, v) in sorted_edges:  
        if not components.in_same_set(u, v):  
            mst.add_edge(u, v)  
            components.union(u, v)
```

```
    # (optional) if mst is now a spanning tree, break  
    if len(mst.edges) == len(graph.nodes) - 1:  
        break
```

```
    return mst
```

Time Complexity

$\Theta(V)$

$\Theta(E \log E)$

```
def kruskal(graph, weights):  
    mst = UndirectedGraph()
```

```
    # place each node in its own disjoint set  
    components = DisjointSetForest(graph.nodes)
```

```
    # sort edges in ascending order by weight  
    sorted_edges = sorted(graph.edges, key=weights)
```

```
    for (u, v) in sorted_edges:  
        if not components.in_same_set(u, v):  
            mst.add_edge(u, v)  
            components.union(u, v)
```

```
    # (optional) if mst is now a spanning tree, break  
    if len(mst.edges) == len(graph.nodes) - 1:  
        break
```

```
    return mst
```

Time Complexity

$\Theta(V)$

$\Theta(E \log E)$

$O(E)$

```
def kruskal(graph, weights):  
    mst = UndirectedGraph()
```

```
    # place each node in its own disjoint set  
    components = DisjointSetForest(graph.nodes)
```

```
    # sort edges in ascending order by weight  
    sorted_edges = sorted(graph.edges, key=weights)
```

```
    for (u, v) in sorted_edges:  
        if not components.in_same_set(u, v):  
            mst.add_edge(u, v)  
            components.union(u, v)
```

```
    # (optional) if mst is now a spanning tree, break  
    if len(mst.edges) == len(graph.nodes) - 1:  
        break
```

```
    return mst
```


Time Complexity

$\Theta(V)$

$\Theta(E \log E)$

$O(E)$

$\alpha(V)$ per line

```
def kruskal(graph, weights):  
    mst = UndirectedGraph()
```

```
    # place each node in its own disjoint set  
    components = DisjointSetForest(graph.nodes)
```

```
    # sort edges in ascending order by weight  
    sorted_edges = sorted(graph.edges, key=weights)
```

```
    for (u, v) in sorted_edges:  
        if not components.in_same_set(u, v):  
            mst.add_edge(u, v)  
            components.union(u, v)
```

```
        # (optional) if mst is now a spanning tree, break  
        if len(mst.edges) == len(graph.nodes) - 1:  
            break
```

```
    return mst
```

Time Complexity

$\Theta(V)$

$\Theta(E \log E)$

$O(E)$

$\alpha(V)$ per line

$O(E \alpha(V))$ for loop

```
def kruskal(graph, weights):  
    mst = UndirectedGraph()
```

```
    # place each node in its own disjoint set  
    components = DisjointSetForest(graph.nodes)
```

```
    # sort edges in ascending order by weight  
    sorted_edges = sorted(graph.edges, key=weights)
```

```
    for (u, v) in sorted_edges:  
        if not components.in_same_set(u, v):  
            mst.add_edge(u, v)  
            components.union(u, v)
```

```
    # (optional) if mst is now a spanning tree, break  
    if len(mst.edges) == len(graph.nodes) - 1:  
        break
```

```
    return mst
```

Time Complexity

$\Theta(V)$

$\Theta(E \log E)$

$O(E)$

$\alpha(V)$ per line

$O(E \alpha(V))$ for loop

```
def kruskal(graph, weights):  
    mst = UndirectedGraph()
```

```
    # place each node in its own disjoint set  
    components = DisjointSetForest(graph.nodes)
```

```
    # sort edges in ascending order by weight  
    sorted_edges = sorted(graph.edges, key=weights)
```

```
    for (u, v) in sorted_edges:  
        if not components.in_same_set(u, v):
```

```
            mst.add_edge(u, v)  
            components.union(u, v)
```

```
    # (optional) if mst is now a spanning tree, break  
    if len(mst.edges) == len(graph.nodes) - 1:  
        break
```

```
    return mst
```

$\Theta(V + E \log E + E \alpha(V))$

Time Complexity

$\Theta(V)$

$\Theta(E \log E)$

$O(E)$

$\alpha(V)$ per line

$O(E \alpha(V))$ for loop

```
def kruskal(graph, weights):  
    mst = UndirectedGraph()
```

```
    # place each node in its own disjoint set  
    components = DisjointSetForest(graph.nodes)
```

```
    # sort edges in ascending order by weight  
    sorted_edges = sorted(graph.edges, key=weights)
```

```
    for (u, v) in sorted_edges:  
        if not components.in_same_set(u, v):  
            mst.add_edge(u, v)  
            components.union(u, v)
```

```
    # (optional) if mst is now a spanning tree, break  
    if len(mst.edges) == len(graph.nodes) - 1:  
        break
```

```
    return mst
```

$\Theta(V + E \log E)$

Time Complexity

$\Theta(V)$

$\Theta(E \log E)$

$O(E)$

$\alpha(V)$ per line

$O(E \alpha(V))$ for loop

```
def kruskal(graph, weights):  
    mst = UndirectedGraph()
```

```
    # place each node in its own disjoint set  
    components = DisjointSetForest(graph.nodes)
```

```
    # sort edges in ascending order by weight  
    sorted_edges = sorted(graph.edges, key=weights)
```

```
    for (u, v) in sorted_edges:  
        if not components.in_same_set(u, v):  
            mst.add_edge(u, v)  
            components.union(u, v)
```

```
    # (optional) if mst is now a spanning tree, break  
    if len(mst.edges) == len(graph.nodes) - 1:  
        break
```

```
    return mst
```

$\Theta(E \log E)$ if assume connected

Time Complexity

- Assume graph is connected. Then $E = \Omega(V)$.
 - $E \geq V - 1$
- Kruskal's takes $\Theta(E \log E) = \Theta(E \log V)$ time.
 - Dominated by sorting the edges.
- **Note:** if graph disconnected, Kruskal's produces a **minimum spanning forest**.

Kruskal vs. Prim

Kruskal v. Prim

- Both algorithms for computing MSTs.
- Which is “better”?
- There’s no clear winner.

Time Complexity

- **Prim:**
 - Binary heap: $\Theta(V \log V + E \log V)$
 - Fibonacci heap: $\Theta(V \log V + E)$
- **Kruskal:** $\Theta(E \log V)$
- If the graph is dense, $E = \Theta(V^2)$, and Prim's with Fibonacci heap "wins".
 - $\Theta(V^2)$ versus $\Theta(V^2 \log V)$.

Not so fast...

- Fibonacci heaps are hard to implement, **high** overhead.
- Prim's will be faster for **very large dense graphs**.
- But Kruskal's may be faster for **smaller dense graphs**.
- The right choice depends on your application.

Main Idea

- Asymptotic time complexity isn't **everything**.
- For small inputs, the “inefficient” algorithm **may** beat the “efficient” one.
- There's also ease of implementation to consider.

Quick check

You run Kruskal's algorithm on a connected, weighted graph with V vertices and E edges.

During execution, you notice that exactly k edges are skipped because they would have formed cycles. Which of the following must be true?

- A. The graph contains exactly k cycles.
- B. The number of connected components in the graph at the start of the algorithm was $k + 1$.
- C. The graph has at least k extra edges beyond what is needed for a spanning tree.
- D. The number of edges with weight ties (equal weights) is at least k .

Did everything make sense about Kruskals?

A: Yes: 100%

B: 80%

C: 60%

D: 40%

E: No :(

MSTs and Clustering

Clustering: Find Groups


Goal: identify the groups in data.

Example:



Clustering, Formalized


- We frame as an **optimization** problem.
- **Given**: n data points.
- **Goal**: assign color to each point (red or blue) to *maximize* the distance between the *closest* pair of red and blue points.



Bad Clustering

Clustering, Formalized

- We frame as an **optimization** problem.
- **Given**: n data points.
- **Goal**: assign color to each point (red or blue) to *maximize* the distance between the *closest* pair of red and blue points.




Good Clustering

Brute Force Solution

- Try all possible assignments; return best.
- If there are n data points, there are $\Theta(2^n)$ assignments.
- **Exponential time**; very slow.
 - Practical only for ~ 50 data points.
- Instead, we will turn it into a **graph problem**.


Distance Graphs

- Given n data points, $p_1, p_2, \dots, p_n$, create *complete* graph with $V = \{p_1, \dots, p_n\}$.
- Set weight of edge $(p_i, p_j) = \text{dist}(p_i, p_j)$.
- The result is a weighted, undirected **distance graph**.





Distance Graphs

- Given n data points, $p_1, p_2, \dots, p_n$, create *complete* graph with $V = \{p_1, \dots, p_n\}$.
- Set weight of edge $(p_i, p_j) = \text{dist}(p_i, p_j)$.
- The result is a weighted, undirected **distance graph**.





Distance Graphs

- Given n data points, $p_1, p_2, \dots, p_n$, create *complete* graph with $V = \{p_1, \dots, p_n\}$.
- Set weight of edge $(p_i, p_j) = \text{dist}(p_i, p_j)$.
- The result is a weighted, undirected **distance graph**.



Distance Graphs



- Given n data points, $p_1, p_2, \dots, p_n$, create *complete* graph with $V = \{p_1, \dots, p_n\}$.
- Set weight of edge $(p_i, p_j) = \text{dist}(p_i, p_j)$.
- The result is a weighted, undirected **distance graph**.



Weight is the distance

Distance Graphs



- Given n data points, $p_1, p_2, \dots, p_n$, create *complete* graph with $V = \{p_1, \dots, p_n\}$.
- Set weight of edge $(p_i, p_j) = \text{dist}(p_i, p_j)$.
- The result is a weighted, undirected **distance graph**.



Weight is the distance

Distance Graphs



- Given n data points, $p_1, p_2, \dots, p_n$, create *complete* graph with $V = \{p_1, \dots, p_n\}$.
- Set weight of edge $(p_i, p_j) = \text{dist}(p_i, p_j)$.
- The result is a weighted, undirected **distance graph**.



Weight is the distance

Distance Graphs



- Given n data points, $p_1, p_2, \dots, p_n$, create *complete* graph with $V = \{p_1, \dots, p_n\}$.
- Set weight of edge $(p_i, p_j) = \text{dist}(p_i, p_j)$.
- The result is a weighted, undirected **distance graph**.



Weight is the distance

Distance Graphs



- Given n data points, $p_1, p_2, \dots, p_n$, create *complete* graph with $V = \{p_1, \dots, p_n\}$.
- Set weight of edge $(p_i, p_j) = \text{dist}(p_i, p_j)$.
- The result is a weighted, undirected **distance graph**.



Weight is the distance

Distance Graphs

- Given n data points, $p_1, p_2, \dots, p_n$, create *complete* graph with $V = \{p_1, \dots, p_n\}$.
- Set weight of edge $(p_i, p_j) = \text{dist}(p_i, p_j)$.
- The result is a weighted, undirected **distance graph**.



You can finish the rest


Main Idea

We can always think of a set of points in a (metric) space as a weighted distance graph.

This is a **very important idea**, because it allows us to use our graph algorithms!

Clustering with MSTs


- Given n data points and a number of clusters, k :
 - Create distance graph G .
 - Run Kruskal's Algorithm on G until there are only k components.



- The resulting connected components are the **clusters**.
- This is known as **single-linkage clustering**.

Clustering with MSTs

- Given n data points and a number of clusters, k :
 - Create distance graph G .
 - Run Kruskal's Algorithm on G until there are only k components.



- The resulting connected components are the **clusters**.
- This is known as **single-linkage clustering**.

Single-Linkage Clustering

- Time complexity of single-linkage is determined by Kruskal's Algorithm: $\Theta(E \log E)$.
- Since distance graph is complete, $E = \Theta(V^2)$, and so $\Theta(E \log E) = \Theta(V^2 \log V) = \Theta(n^2 \log n)$
- Practically, can cluster ~ 10, 000 points.

Summary

- We started the quarter with a **brute force solution**.
 - Took $\Theta(2^n)$ time, only feasible for a few dozen points.
- We've now reframed the problem using graph theory.
 - Now only $\Theta(n^2 \log n)$ time!
 - Feasible for tens of thousands of points.

Why Algorithms?

- Data scientists use computers as tools.
- But solving a problem isn't just about coding it up.
- You need to know how to analyze your code and use the right algorithms and data structures to make your solution efficient.



Thank you!

Do you have any questions?

CampusWire!